

Solving Mazes: Dyna-Q+ Agent

[Achi's Projects \(https://github.com/QuantumNano-AI/PROJECTS\)](https://github.com/QuantumNano-AI/PROJECTS)

The DynaQ+ algorithm is more robust than DynaQ. It provides a solution to environment dynamics changing and thus making its model inaccurate

For the DynaQ+ Agent we will be adding two new parameters state visitation count τ and scaling parameter κ .

Where r is the modeled reward for a transition which hasn't been visited in $\tau(s, a)$ time steps, the planning updates are done using $r + \kappa\sqrt{\tau(s, a)}$ as the reward instead


```

        reward = self.rewards['goal'] if ((r,c) == self.goal) or (state == self.goal) else self.rewards['other']
        trans_prob = 1
        result.append((action, (r, c), reward, trans_prob, terminal))

    actions = [tup[0] for tup in result]

    if a:
        R = []
        if a in actions:
            inx = actions.index(a)
            R.append((result[inx]))
            return R
        else:
            R.append((a, (row,col), self.rewards['wall'], 1, terminal))
            return R
    return result

def plot_state_values(self):
    val = np.array(list(self.V.values())).reshape(-1,1)
    va = sc.MinMaxScaler(feature_range=(0, 255)).fit_transform(val).flatten()
    V = {}
    for i in range(len(va)):
        V[list(self.V.keys())[i]] = va[i]

    # create a black image
    img = np.ones((self.height,self.width,3), np.uint8)

    for item in V.items():
        (r,c),vx = item
        img[r,c] = [0,vx,0]

    for r,c in self.wall_cords:
        img[r,c] = [150,5,150]

    img[self.start[0],self.start[1]] = [255,0,0]
    img[self.goal[0],self.goal[1]] = [100,100,255]
    def showing(img):
        plt.figure(figsize = (15,15))
        plt.imshow(img, cmap='viridis')
        plt.xticks([])
        plt.yticks([])
        #plt.colorbar()
        plt.show()
    showing(img)

    def policy_(s):
        row, col = s
        candidates = [
            ("up", (row - 1, col)),
            ("down", (row + 1, col)),
            ("left", (row, col - 1)),
            ("right", (row, col + 1))
        ]

        if s in self.wall_cords:
            return ('WALL!!!')
        else:
            values = {a:self.V[r,c] for a,(r,c) in candidates if 0 <= r < self.height and 0 <= c < self.width and not self.walls[r][c]}
            values = {v:k for k,v in values.items()}
            best = values[max(values)]
            return best

    pi = np.zeros((self.height, self.width)).astype('str')
    pi = np.where(pi=='0.0', 'wall', pi)
    candidates = {
        "up": [255,0,0],
        "down": [0,0,255],
        "left": [0,255,0],
        "right": [255,255,0]
    }
    for item in self.pi.keys():
        action = policy_(item)

        r,c = item
        pi[r,c] = action
        img[r,c] = candidates[action]

    img[self.start[0],self.start[1]] = [255,0,0]
    img[self.goal[0],self.goal[1]] = [100,100,255]
    def showing(img):
        plt.figure(figsize = (15,15))
        plt.imshow(img, cmap='viridis', )
        plt.yticks(list(range(self.height)))
        plt.xticks(list(range(self.width)))
        plt.title("POLICY\n\nRED => up\nBLUE => down\nGREEN => left\nYELLOW => right")
        #plt.colorbar()
        plt.show()
    showing(img)
    return img

```

```

In [2]: class DynaQ(Maze):
def __init__(self, maze, rewards = {'goal':1000000,
                                     'wall':-15,
                                     'other':-1},
             info = {} ):

    Maze.__init__(self, maze, rewards)
    self.epsilon = info.get('epsilon', 0.9) # Exporation parameter
    self.r = np.random.RandomState(seed=12345)
    self.episodes =info.get('episodes', 200)
    self.max_steps =info.get('max_steps', 1500)
    self.alpha = info.get('alpha',0.1) # step size
    self.gamma = info.get('gamma',.99) # discount factor
    self.kappa = info.get('kappa',.001) # Scalling parameter
    self.planning_steps = info.get('planning_steps',50)
    self.tau = pd.DataFrame(np.zeros((self.state_count, len(self.actions))),columns=self.actions, index=self.states).to_dict(orient='index')

def func_q(self, states,n_states,n_actions,kind = 'random'):
    if kind=='ones':
        return dict(zip(states,np.ones((n_states,n_actions)).tolist()))
    elif kind == 'zeros':
        return dict(zip(states,np.zeros((n_states,n_actions)).tolist()))
    elif kind == 'random':
        return dict(zip(states,np.round(self.r.randn(n_states,n_actions),2).tolist()))
    else : raise NameError("Wrong input: please use ['ones', 'zeros', 'random']")

def argmax(self, test_array):
    return self.r.choice(np.flatnonzero(np.array(test_array)==np.array(test_array).max()))

def epsilon_greedy(self, Q, epsilon, actions, state, train=False):
    current_q = Q[state]
    if self.r.rand() < epsilon:
        action = self.r.choice(actions)
        return action
    else:
        action = self.argmax(current_q)
    return actions[action]

def sigmoid(self,a):
    import numpy as np
    s = np.divide(1,1+np.exp(-a))
    return s

def run(self):
    self.Q = self.func_q(self.states,self.state_count,len(self.actions),kind = 'zeros')
    self.model = {} # model is a dictionary of dictionaries, which maps states to actions to (reward, next_state) tuples
    def update_model(s,a,s_,reward):
        if s in self.model: self.model[s][a] = (s_,reward) # If the agent has been in this state before, update the action/reward
        else:
            self.model[s] = {a:(s_,reward)} # else add new state and action to model
            # actions that had never been tried are now going to be considered. The initial model for such actions would be
            # to lead back to the previous state and have a reward of zero
            for action in self.actions:
                if action != a:
                    self.model[s][a] = (s, 0)

    def planning():
        for i in range(self.planning_steps):
            s = list(self.model.keys())[self.r.randint(len(self.model.keys()))]
            a = self.r.choice(list(self.model[s].keys()))
            (s_,reward) = self.model[s][a]
            reward += self.kappa*np.sqrt(self.tau[s][a])
            q_ = self.Q[s_]
            a_ = self.actions[self.argmax(q_)] # Action in the next state does not follow policy. It is rather selected to maximise utility
            if terminal:
                self.Q[s][self.actions.index(a)] += self.alpha * (reward - self.Q[s][self.actions.index(a)])
            else:
                self.Q[s][self.actions.index(a)] += self.alpha * (reward + self.gamma*max(q_) \
                    - self.Q[s][self.actions.index(a)])

    self.time_steps = pd.DataFrame()

    for episode in range(self.episodes):

        total_reward = 0 # This sets the total reward obtained during this episode
        s = self.states[self.r.randint(len(self.states))]
        a = self.epsilon_greedy(Q=self.Q, epsilon=self.epsilon, actions=self.actions, state=s)
        t = 0
        terminal = False
        while t < self.max_steps:
            t+=1
            self.tau = {k:{kk:vv+1 for kk,vv in v.items()} for k,v in self.tau.items()}
            self.tau[s][a] = 0

            _,s_, reward, p,terminal = self.neighbors(s,a)[0]
            total_reward += reward
            q_ = self.Q[s_] # Action values in the next state
            a_ = self.epsilon_greedy(Q=self.Q, epsilon=self.epsilon, actions=self.actions, state=s)
            reward += self.kappa*np.sqrt(self.tau[s][a])
            if terminal:
                self.Q[s][self.actions.index(a)] += self.alpha * (reward - self.Q[s][self.actions.index(a)])
            else:
                self.Q[s][self.actions.index(a)] += self.alpha * (reward + self.gamma*self.Q[s_][self.actions.index(a_)] \
                    - self.Q[s][self.actions.index(a)])

            update_model(s,a,s_,reward)

        # Carry out planning only when there is a complete episode with rewards returned
        if len(self.time_steps)>0:
            ts = self.time_steps[self.time_steps.rewards!=0]
            if len(ts)>0:planning()

```

```

s, a = s_, a_
if terminal:
    self.time_steps = self.time_steps.append(
        pd.Series({'episode':int(episode), 'steps':t, 'rewards':total_reward}), ignore_index=True)
    break
if t%10==0:
    print(f'.',end='')
self.pi = {}
self.V = {}
for k,v in self.Q.items():
    self.pi[(k)] = self.actions[self.argmax(v)]
    self.V[(k)] = max(v)
max_r = max([v for k,v in self.V.items()])
self.V[self.goal] = max_r
img = self.plot_state_values()
#return V,pi

```

```

In [3]: maze0 = """#####
A      ##
## ##### # #
## ##### # #
##      # #
## ##### #
##      ##### #
#####B#"""

```

```

In [4]: print(maze0)

```

```

#####
A      ##
## ##### # #
## ##### # #
##      # #
## ##### #
##      ##### #
#####B#


```

```

In [5]: info = {'episodes': 50, 'max_steps': 5000, 'alpha': 0.1, 'epsilon':.1, 'planning_steps':20, 'kappa':.5}
rewards = {'goal':10000, 'wall':-1, 'other':0}
m = DynaQ(maze = maze0, rewards = rewards, info = info)
m.run()
#t = m.time_steps; t[t.rewards!=0][['episode', 'steps']]

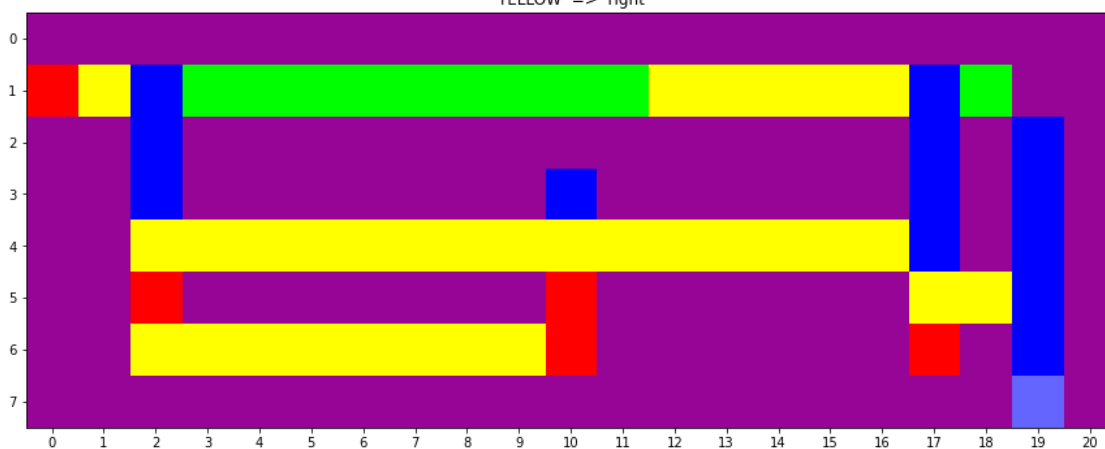
```

.....



POLICY

RED => up
 BLUE => down
 GREEN => left
 YELLOW => right



In []:

In []:


```
In [6]: maze = """#####
##      ##
##B### #####  ##
###          ##
#   #####  ##
#   #####      ##
#           # #####
# ##### # #####
#   ##  ##
#   #####  ##
##  ##      ##
## #####  ##
## #####  ##
##      ##
###A#####"""
```

```
In [7]: print(maze)

#####
##      ##
##B### #####  ##
###          ##
#   #####  ##
#   #####      ##
#           # #####
# ##### # #####
#   ##  ##
#   #####  ##
##  ##      ##
## #####  ##
## #####  ##
##      ##
###A#####
```


```
In [ ]:
```

```
In [8]: info = {'episodes': 200, 'max_steps': 5000, 'alpha': 0.1, 'epsilon':.1, 'planning_steps':100, 'kappa':.5}
rewards = {'goal':10000000, 'wall':-50, 'other':0}
m = DynaQ(maze = maze, rewards = rewards, info = info)
m.run()
#t = m.time_steps; t[t.rewards!=0][['episode', 'steps']]
.....
```



POLICY

RED => up
 BLUE => down
 GREEN => left
 YELLOW => right



In []:

In []:

```
In [9]: maze1 = """#####
##          #
## ### ##### #
##          ### #
## ### #####B#
##### #####
##          ##### #
## ### # ##### #
## ###      #
## #####      ### #
## #          ##### ### #
## # #####      ### #
## # #####      ### #
## # #####      ### #
##          #
## #####      #
##A#####"""
```

In [10]:


```
print(maze1)

#####
##          #
## ### ##### #
##          ### #
## ### #####B#
##### #####
##          ##### #
## ### # ##### #
## ###      #
## #####      ### #
## #          ##### ### #
## # #####      ### #
## # #####      ### #
## # #####      ### #
##          #
## #####      #
##A#####
```



```
In [11]: info = {'episodes': 50, 'max_steps': 5000, 'alpha': 0.1, 'epsilon':.1, 'planning_steps':20}
rewards = {'goal':100000000, 'wall':-100, 'other':0}
m = DynaQ(maze = maze1, rewards = rewards, info = info)
m.run()
#t = m.time_steps; t[t.rewards!=0][['episode', 'steps']]
```

....



POLICY
RED => up
BLUE => down
GREEN => left
YELLOW => right

